XYZ-guard — Positioning & Messaging One-Pager

For the open-source launch. Internal working document.

Positioning Statement

For platform-security and AI-governance teams **putting AI agents into production**, **who** cannot risk an agent taking a catastrophic action (dropping a table, spamming customers, burning budget in a loop), **XYZ-guard is** a runtime action-safety layer **that** scores every tool call an agent makes on a 0–100 risk scale and routes medium-risk actions to a human for approval — with built-in behavioral safety and zero agent code changes. **Unlike** LLM gateways (LiteLLM, Portkey, Bifrost, Kong) that route and meter model traffic but make binary allow/deny tool decisions, and **unlike** content guardrails (Lakera, NeuralTrust) that screen what an agent *says*, **XYZ-guard governs what an agent *does*** — at the tool call, with a graded human-in-the-loop tier no competitor offers.

The One-Liner (elevator)

XYZ-guard is the action firewall for AI agents — it scores every tool call and puts a human in the loop before anything risky happens, in any language, with zero safety code.

Category

Agent action-safety / runtime guardrail — deliberately *not* "AI gateway." The gateway category is crowded and consolidating; the action-safety niche is open, and the guardrails market is growing faster (~32.5% vs ~29.4% CAGR).

Three Messaging Pillars

Pillar	Headline	Proof point
1. Graded enforcement	"Not just allow or block — a human-review tier for the risky middle."	0–100 risk score → allow (≤30) / intervene (31–70) / block (>70). No competitor has a scored intervene tier.
2. Behavioral safety built in	"Catches runaway agents, not just bad prompts."	Loop / drift / hallucination / contradiction detection + circuit breaker + checkpoint-rollback,

		bundled.
3. Zero-code, any language	"Point your agent at a URL. Done."	Three deployment modes from one image, incl. transparent NAT — agent doesn't even know it's there. Deterministic, in-process, no LLM in the eval path, no phone-home.

Messaging by Audience

- **Platform-security / governance buyer (primary):** "Centralized, un-bypassable policy on every agent action, with an audit trail and a human-approval tier — the control plane your security team needs before agents go to prod."
- **Platform engineer deploying agents:** "Stop building your own safety checks into every agent. One gateway, any framework, policy hot-reloads without a restart."
- **Individual agent developer:** "Five lines, or zero — you don't even have to know XYZ-guard is there. Your agent just can't do the catastrophic thing."

Proof Points to Lead With

- **Framework-agnostic:** OpenAI, Claude, Bedrock, Strands (+ any via a ~50-line adapter).
- **Deterministic & private:** no LLM calls during evaluation; runs in-process; no telemetry / phone-home.
- **Production signals:** 585 tests, 90%+ coverage, strict typing, CI across Python 3.10–3.13; Apache-2.0.
- **Complete platform:** core library + Agent Gateway + 12-capability control plane (registry, policies, quotas, cost, A/B, live traces, audit).

Competitive Contrast (say this, not that)

Say	Don't say
"Action-level safety with a human-review tier."	"Another AI gateway."
"Governs what agents *do* ."	"Routes LLM traffic across providers." (that's table stakes)
"Per- *call* risk scoring."	"Risk scoring" alone (Lasso Security scores

	agents — collision risk).
"Catches runaway agents (loops, drift)."	"Content safety / prompt-injection." (that's Lakera's lane)

Objection Handling

- **"Isn't this just a gateway?"** → No — gateways route and meter; XYZ-guard scores and *intervenes* on actions. Routing is one feature, not the point.
- **"How is it different from Lakera / guardrails?"** → They screen text (what the agent says). XYZ-guard governs tool calls (what the agent does). Complementary, not competing.
- **"You're Beta / only 6 providers."** → Provider breadth isn't the value; action-safety is. Bring your own gateway if you want 200 providers — XYZ-guard is the safety layer on top.
- **"Why should I trust an OSS safety tool after the LiteLLM compromise?"** → XYZ-guard is in-process, deterministic, makes no external calls during evaluation, and phones home to no one. Auditable by design.

Naming Watch-Out

"**Lasso Security**" ships an MCP security gateway and uses "risk score" language (for *agents*, not calls). Expect confusion. Mitigate in every asset: always pair the name with "**per-call action safety**," and never use bare "risk score" — always "**per-tool-call risk score**."

Launch Headline Options

1. *"XYZ-guard: the action firewall for AI agents."*
2. *"Score every tool call. Keep a human in the loop. In any language."*
3. *"The guardrail layer between your agents and the outside world."* (from the repo — strong, keep as tagline)

What NOT to Lead With

Provider breadth · routing/failover · cost tracking · MCP support. These are **table stakes** — mention them as "yes, we do that too," never as the pitch. Leading with them drops XYZ-guard into the crowded gateway race it cannot win on funding.