

XYZ-guard: Comparative Market Analysis & Competitor Market Share

Comprehensive competitive analysis using the full XYZ-guard platform (Core Library + Agent Gateway/Sidecar + Control Plane) as the baseline. Feature claims are grounded in the XYZ-guard repository; competitor and market data carry inline caveats where vendor-sourced.

1. Executive Summary

XYZ-guard is a three-layer, safety-first runtime platform for AI agents: a **Core Library** (in-process Guard with policy engine, four anomaly detectors, circuit breaker, checkpoint/rollback), an **Agent Gateway** (language-agnostic proxy with a two-path model — direct `/tool/{action}` and the LLM-driven `/invoke` agentic loop, powered by the ABC-LLM routing engine, dual authorization, dual quota, and a Tool Router spanning HTTP/MCP/Agent-as-Tool), and a **Control Plane** (12-capability web console pushing config to a gateway fleet and collecting traces, cost, and audit).

Assessed against this fuller baseline, the conclusion is unchanged but sharper: **no single competitor replicates XYZ-guard's full stack**, yet on nearly every individual capability a stronger, better-resourced competitor exists. XYZ-guard's defensible wedge is the combination of (a) a per-call **0–100 risk score → allow/intervene/block** with an automated human-in-the-loop *intervene* tier, (b) bundled **behavioral-safety** depth (loop/drift/hallucination/contradiction detection), and (c) **three deployment modes from one image** including zero-agent-code transparent NAT — delivered on an Apache-2.0 open core.

On **market share**, the essential finding is that **no analyst publishes product-level share** for this ~18-month-old, open-source-dominated category. Share must be inferred from proxies: GitHub mindshare (LiteLLM leads by a wide margin), strategic consolidation (Palo Alto Networks acquiring Portkey), and adoption signals. XYZ-guard itself holds **no measurable share** today — it is a Beta entrant not yet indexed in any comparison.

2. The Competitive Set

XYZ-guard sits where two waves collide on the same tool-call chokepoint: AI/LLM gateways extending downstream into MCP tool governance, and AI-security/guardrail platforms extending into the tool layer. The main competitors, by category:

- **AI/LLM & MCP gateways:** agentgateway.dev (Solo.io → Linux Foundation), Bifrost (Maxim AI), LiteLLM (BerriAI), Portkey (→ Palo Alto Networks), Kong AI Gateway, Cloudflare AI Gateway, ToolHive (Stacklok/Red Hat), Lunar MCPX.
- **Agent runtime governance / guardrails:** Lakera Guard (Check Point), NeuralTrust, Lasso Security MCP Gateway, Invariant Labs.

The closest head-to-head competitors to XYZ-guard's *runtime enforcement* core are **agentgateway** (architecture), **Bifrost** (LLM+MCP+cost bundle), **Lakera Guard** (runtime tool-call control), and **Lasso Security MCP Gateway** (security intent + name-confusion risk).

3. Comparative Capability Analysis

Legend: ✓ strong/native · ● partial/adjacent · ✕ absent. XYZ-guard is the reference frame.

3a. Governance & Safety Enforcement

Dimension	XYZ-guard	agentgateway	Bifrost	Lakera Guard	Lasso Sec. MCP GW	ToolHive	Lunar MCPX
Tool-call policy enforcement	✓ every call (Guard)	✓ CEL RBAC per tool	✓ per virtual key	✓ allow/deny at invocation	✓ proxy governance MCP	✓ per-request Cedar	✓ per-agent scope
Risk scoring (0–100 numeric)	✓ per call	✕	✕	● sensitivity	● per-agent	✕	✕
Three-tier allow/intervene/block	✓ ≤30/≤70/>70	● binary	● suggest + allow-list	● flag/block	● block high-risk agent	● allow/deny/modify	● allow/deny
Human-in-the-loop intervene tier	✓ intervention gateway	✕	● suggestion mode	✕	✕	● webhook	✕
Anomaly detection (loop/drift/hallucination/contradiction)	✓ 4 detectors	✕	● additions	✓ prompt-attack	✓ behavioral	✕	✕

Circuit breaker / checkpoint-rollback	✓ both	×	×	×	×	×	×
Bypass-proof (no direct agent→tool path)	✓ archite ctural	✓ credenti al isolatio n	✓ crede ntial isolat ion	● inline scree ning	✓ proxy	✓ container isolation	✓ middl eware

3b. Protocols, Routing & Cost

Dimension	XYZ-guard	agentgat eway	Bifrost	Lak era Gua rd	Las so Sec. MC P GW	ToolH ive	Luna r MCP X
MCP support	✓ embedded/remo te/stdio	✓ federation + OpenAPI →MCP	✓ client+s erver	● scre ens calls	✓ MC P- nati ve	✓ lifecyc le + proxy	✓ multi - serve r
Agentic loop (intent→plan→s ynthesis)	✓ /invoke + ABC-LLM	● proxy only	● agent mode	×	×	×	● dyna mic select ion
A2A support	×	✓ native	● agents framing	×	×	×	×
LLM multi- provider routing	✓ 6 providers	✓ ~20	✓ 23+ / 1000+ models	×	×	×	×
Quota / budget / cost (runtime enforced)	✓ dual, pre-LLM projection	✓ denial- of-wallet	✓ hierarch ical budgets	×	×	×	● metri cs
Agent-as-Tool	✓	● via A2A	●	×	×	×	×

routing

3c. Deployment, Control Plane, Observability & Maturity

Dimension	XYZ-guard	agentgateway	Bifrost	Lakera Guard	Lasso Sec. MCP GW	ToolHive	Lunar MCPX
Deployment modes	✓ sidecar/shared/NAT (1 image)	✓ VM/container/K8s	✓ single binary, in-VPC	● managed API	● self-host	● local + K8s	● self-host
Zero-agent-code transparent mode	✓ Istio/Envoy/DNS	● URL repoint	● drop-in	✓ no prompt changes	● compatible	● proxy	● middleware
Control plane & hot-reload	✓ 12-capability console	✓ xDS controller	✓ web UI	● policy console	● config-driven	✓ K8s operator	✓ central control
Fleet mgmt (registry/AB/sandbox/audit)	✓ full	● partial	● partial	●	✗	●	●
Observability / tracing	✓ OTel + live WS viewer	✓ OTel	✓ Prometheus + OTel	✓ logs	✓ audit	✓ platform	✓ visibility
Agent auth	● deny-by-default (JWT planned)	✓ JWT/OAuth OBO	✓ virtual keys + SSO	● detector	● plugin	✓ OIDC + Cedar	✓ per-agent ACL
Licensing / backing	Apache-2.0, indie	Apache-2.0, Linux Foundation	Apache-2.0, Maxim AI	Commercial, Check Point	MIT, Lasso Security	Apache-2.0, Stacklok/Red Hat	OSS + commercial

Maturity	Beta (585 tests)	v1.2–1.3 + enterprise	Vendor-benchmarked	Production	Launched May 2026	Active, K8s lineage	Production
----------	------------------	-----------------------	--------------------	------------	-------------------	---------------------	------------

Reading the matrix: XYZ-guard is uniquely strong on the *safety* rows (risk scoring, intervene tier, anomaly detection, circuit breaker/rollback) and the *deployment-mode* breadth, and now — with the full platform — also matches the field on control-plane/fleet management and the agentic loop. It remains behind on provider breadth, shipped auth, enterprise backing, and maturity.

4. Market Share of the Competitors

Critical caveat up front: there is **no analyst-published, product-level market share** (revenue or unit %) for AI gateways. Gartner tracks the category, not vendor splits [Gartner — Market Overview for AI Gateways](#). The category is ~18 months old and dominated by open-source, self-hosted tools that never report usage. "Share" below is therefore built from the best available proxies, each with limits noted.

4a. Market size (the envelope)

Segment	2025/26 size	2030 forecast	CAGR
AI inference gateways	\$2.71B (2025) → \$3.5B (2026)	\$9.83B	~29.4%
LLM inference guardrails	\$1.96B (2025)	\$7.99B	~32.5%
LLM middleware gateway (penetration)	~42% of enterprises	—	~49.6% (to 2034)

Sources: Business Research Company (AI Inference Gateways; LLM Inference Guardrails reports); Maxim (vendor-sourced penetration/CAGR). Treat single-vendor figures as directional.

<image: market_size.png>

Figure 1: The two markets XYZ-guard straddles — gateways (~\$3.5B in 2026) and guardrails (~\$2B), both growing ~30%+ annually.

4b. Proxy #1 — Open-source mindshare (GitHub stars)

The least-bad share proxy for an OSS-heavy market:

Product	GitHub stars (approx.)	Providers	Backing / posture
LiteLLM	~52,600	100+	BerriAI (YC, ~\$2.1M raised, ~13 staff)
Portkey	~12,100	200+	Being acquired by Palo Alto Networks
Bifrost	~6,260	23+ / 1000+ models	Maxim AI
agentgateway	growing (LF-backed)	~20	Linux Foundation (AWS/Cisco/IBM/MS/Red Hat)
XYZ-guard	early / Beta	6	Indie, Apache-2.0

Interpretation: LiteLLM is the runaway mindshare leader — roughly **4× Portkey** and **8× Bifrost** — and is the product XYZ-guard is most often benchmarked against. Provider breadth ranks Portkey (200+) > LiteLLM (100+) > Kong (30+) > Cloudflare (~10). Caveat: stars measure interest, not deployments or revenue.

<image: mindshare.png>

Figure 2: Open-source mindshare (GitHub stars, 2026). LiteLLM leads by ~4x over Portkey and ~8x over Bifrost. Stars measure interest, not revenue or deployments.

4c. Proxy #2 — Strategic consolidation (enterprise share signals)

- **Portkey → Palo Alto Networks** — the single biggest move; Portkey (which processes 1T+ tokens/day across thousands of teams) becomes the AI Gateway for Prisma AIRS. This locks up a large slice of *managed-gateway* enterprise share behind a security major.
- **agentgateway → Linux Foundation** — strongest institutional/vendor-neutral backing.
- **ToolHive** — built by the team behind Kubernetes, maintained by Stacklok + Red Hat.

- **Kong** — first API platform to ship a dedicated AI gateway (Feb 2024), riding its installed API-management base.

4d. Proxy #3 — Deployment-pattern signal (qualitative)

Across enterprise conversations and 2026 buyer's guides, the dominant patterns are consistent: most enterprises either **adopt LiteLLM (self-hosted)**, **evaluate Kong or Portkey** for managed/enterprise governance, or **build a custom DIY gateway** — with DIY common among the most strategic customers. LiteLLM's true footprint is systematically understated by any measurable signal because the majority of its users self-deploy via Docker/GitHub and are never counted. This qualitative pattern — not a precise count — is the most defensible adoption read available.

4e. Estimated share of position (synthesis)

Combining the proxies into a qualitative ranking (there is no defensible precise percentage):

Tier	Products	Basis
Mindshare / volume leader	LiteLLM	~52.6K GitHub stars; most-named in 2026 buyer's guides; dominant self-hosted OSS
Managed / enterprise leader	Portkey (now Palo Alto), Kong	1T+ tokens/day; enterprise API base
Performance challenger	Bifrost	Go perf claims; 23+ providers
Agent-native / protocol	agentgateway	MCP + A2A, LF backing
Guardrails / security-first	Lakera, NeuralTrust, Lasso Security	runtime screening; guardrails market ~\$2B
DIY	custom-built	plurality of strategic enterprises
XYZ-guard	—	~0% measured share — pre-index Beta

<image: share_tiers.png>

Figure 3: Estimated share of position — a qualitative, proxy-based ranking (no defensible precise percentages exist for this category).

4f. Competitor Customer Counts (disclosed / detected)

Customer counts are only knowable where a vendor publishes them or a third party detects deployments. The managed/commercial vendors disclose numbers; pure open-source projects mostly do not, because users self-deploy invisibly. The figures below are **not apples-to-apples** — Portkey's is self-reported org signups, while the Kong/LiteLLM figures are third-party *detected* deployments (a floor).

Competitor	Customer figure	What it measures
Portkey	24,000+ organizations globally (also cites "3,000+ teams")	Self-reported orgs on the platform
Kong (overall)	~1,430 companies (~139 on Kong Enterprise)	Third-party detected deployments
LiteLLM	~827 companies detected	Third-party detected deployments (floor; most self-host invisibly)
Bifrost (Maxim AI)	Not disclosed	Publishes performance/feature claims, not customer counts
agentgateway	Not disclosed	Linux Foundation project — contributor orgs named, not customers
Lakera / NeuralTrust / Lasso	Not disclosed as counts	Named enterprise logos only

Portkey is the standout: it publicly reports processing "more than 500B LLM tokens across 125M requests per day, and managing more than \$500,000 in AI spend daily for more than 24,000 organizations globally" [Portkey Series A — GlobeNewswire](#){evidence: "managing more than \$500,000 in AI spend daily for more than 24,000 organizations globally"} and describes itself as "trusted by 3000+ teams, from Fortune 500s to startups" [portkey.ai](#){evidence: "Portkey is trusted by 3000+ teams, from Fortune 500s to startups, processing billions of tokens daily."}. This is the most credible figure in the set because it comes from a funding announcement. Third-party scans detect "1,430 companies using Kong" [bloomberly — Kong](#){evidence: "We detected 1,430 companies using Kong."} and "827 companies using LiteLLM" [bloomberly — LiteLLM](#){evidence: "We detected 827 companies using LiteLLM. The most common industry is Software Development (21%)"},

but these badly under-count self-hosted installs — LiteLLM's packages alone see roughly 3M downloads/day. Named enterprise customers where disclosed: LiteLLM cites Rocket Money, Samsara, Lemonade, and Adobe; Portkey cites Postman, Snorkel AI, and Fortune 500s across finance, pharma, and technology.

Bottom line: only Portkey publishes a true customer number (24,000+ orgs) and leads on disclosed adoption — reinforcing the Palo Alto acquisition rationale. Kong (~1,430) and LiteLLM (~827) have third-party detected floors, not official figures. Everyone else keeps counts private, so a clean side-by-side customer-count ranking across the field is not possible.

4g. Competitor Financials (revenue, funding, valuation, M&A)

Almost none of these companies disclose true revenue — most are private or early-stage. Only **Kong** publishes actual ARR. For the rest, the best public figures are **funding raised**, **acquisition price**, or **annualized spend managed** (a usage proxy, not the vendor's own revenue). These are **not comparable as "revenue."**

Competitor	Financial figure	Type
Kong (parent of Kong AI Gateway)	\$100M+ ARR (est. ~\$146M); \$2B valuation; raised \$175M Series E	Actual ARR + valuation
Portkey (→ Palo Alto)	\$15M Series A (Elevation, Lightspeed); manages \$180M+ annualized LLM spend (customers', not Portkey revenue)	Funding + usage proxy + acquisition
Lakera (→ Check Point)	Acquired for ~\$300M (Sept 2025)	Acquisition price
NeuralTrust	\$20M seed (June 2026 — largest EU cybersecurity seed to date)	Funding
BerriAI (LiteLLM)	~\$1.6–2.1M seed (Y Combinator, Gravity, Pioneer); ~13 employees	Funding
Maxim AI (Bifrost)	Not disclosed	—
agentgateway	N/A — Linux Foundation project, not a company	—

Kong is the only competitor with disclosed revenue, having announced "surpassing \$100 million in Annual Recurring Revenue (ARR)" [Kong — \\$100M ARR](#){evidence: "surpassing \$100 million in Annual Recurring Revenue (ARR)"} at a \$2B valuation on a \$175M Series E [Kong Series E](#){evidence: "\$175 million in up-round Series E financing, with a mix of primary and secondary transactions at a \$2 billion valuation"} — but that figure spans Kong's entire API+AI business, of which the AI gateway is only a slice. Portkey raised "\$15 million in Series A funding led by Elevation Capital, with participation from Lightspeed" [Portkey Series A](#){evidence: "raised \$15 million in Series A funding led by Elevation Capital, with participation from Lightspeed"}; its widely-cited \$180M+ figure is annualized LLM spend its customers run through the platform, not Portkey's own revenue. Lakera was acquired by Check Point "in a deal valued at \$300 million" [ET Telecom](#){evidence: "acquiring Lakera... in a deal valued at \$300 million"}, and NeuralTrust raised "a \$20 million seed round, the largest cybersecurity seed financing raised by an EU company to date" [NeuralTrust seed](#){evidence: "a \$20 million seed round, the largest cybersecurity seed financing raised by an EU company to date"}. LiteLLM's parent BerriAI raised only "a \$1.6M seed round from Y Combinator, Gravity Fund and Pioneer" [YC — LiteLLM](#){evidence: "We have raised a \$1.6M seed round from Y Combinator, Gravity Fund and Pioneer"}, reflecting an early-stage OSS company whose revenue is likely small despite very large usage.

Bottom line: there is no clean revenue ranking for this field. The only actual revenue figure is Kong's \$100M+ ARR (across its whole API business, not just the AI gateway). Everything else is funding, acquisition price, or a usage proxy. The clearest financial signal is not revenue but the consolidation wave — Lakera → Check Point (~\$300M) and Portkey → Palo Alto — which shows where strategic value is being placed.

5. Positioning & Strategic Read

- **The competitor to beat is LiteLLM and "build-your-own,"** not agentgateway or Bifrost — those are strong on paper but show little named enterprise adoption in internal data, whereas LiteLLM and DIY dominate actual deployments.
- **The market is consolidating fast** (Palo Alto→Portkey, LF→agentgateway, Stacklok/Red Hat→ToolHive), raising the bar for a standalone indie entrant.
- **XYZ-guard's fuller platform closes the "table-stakes" gap** — with the control plane, agentic / *invoke* loop, ABC-LLM routing, and dual quota, it now matches the field on fleet management and routing, so the conversation can shift to its safety differentiators rather than feature parity.
- **The defensible niche is the gateway/guardrails seam:** the fastest-growing segment (guardrails, ~32.5% CAGR) meets the enforcement chokepoint, and no gateway leader offers per-call risk-scored allow/intervene/block with behavioral safety. That is the wedge that also explains why so many customers currently *build their own*.
- **Recommended play:** lead with the risk-scored intervene tier + behavioral-safety layer + zero-code deployment; treat routing/cost/MCP as parity table-stakes; and close the credibility gaps that most threaten enterprise adoption — published

performance benchmarks, shipped JWT/identity auth, and compliance attestations (SOC 2 / HIPAA / ISO), which competitors like Bifrost already carry.

6. Go-to-Market & Niche Strategy (Pre-Open-Source)

6.1 The market you are entering

XYZ-guard is entering a **consolidating**, not empty, market. In recent months Palo Alto Networks acquired Portkey, Check Point acquired Lakera (~\$300M), and agentgateway moved to the Linux Foundation with AWS/Cisco/IBM/Microsoft/Red Hat behind it. Two implications:

- **Do not compete on the crowded axis** — provider breadth, routing, cost tracking. LiteLLM, Portkey, Bifrost, and Kong already own it and are better funded; a feature-parity race is unwinnable.
- **Consolidation is happening in security, not gateways.** Both major acquisitions were security vendors buying agent-safety capability — the clearest signal of where the value (and budget) sits.

6.2 The niche to own — the "agent action firewall"

Doing more than competitors is not a position. The niche is the one thing none of them do, stated sharply:

XYZ-guard is the runtime action-safety layer for AI agents — the only gateway that scores every tool call 0–100 and routes medium-risk actions to a human, with built-in behavioral safety (loop / drift / hallucination / contradiction) — dropped in with zero agent code.

Three defensible pillars, in priority order:

1. **Risk-scored intervene tier** — every competitor is binary allow/deny; XYZ-guard has graded allow / human-review / block. Genuinely unique.
2. **Behavioral-safety depth** — anomaly detectors + circuit breaker + checkpoint/rollback bundled in. Gateways don't carry this; guardrail vendors do content, not actions.
3. **Zero-agent-code deployment** across three modes (including transparent NAT).

MCP, routing, cost, and the control plane are **table stakes XYZ-guard has reached** — cited to prove completeness, not led with.

6.3 The market to target

- **Buyer:** the platform-security / AI-governance owner, not the individual agent developer. The consolidation proves budget flows to security buyers; the control plane + audit log + intervene tier is a security story.
- **Beachhead — the "build-your-own" segment.** The plurality of strategic enterprises build custom gateways precisely because nothing off-the-shelf does

action-safety their way. These teams are already convinced they need this — the shortest path to adoption.

- **Wedge use case:** high-blast-radius agents — DevOps/infra, financial ops, customer-data actions — where a wrong tool call is expensive and a human-review tier is worth paying for. Regulated industries (finance, pharma) pushed by the EU AI Act's August 2026 high-risk enforcement toward exactly this kind of runtime oversight.

6.4 Open-source readiness (beyond the code)

- **Name-collision risk — resolve before the repo is public.** "Lasso Security" already ships an MCP security gateway and uses "risk score" language (for agents, not calls). Differentiate hard on messaging ("action-level, per-call") or reconsider naming now.
- **Credibility gaps competitors will exploit:** Beta, Python, 6 providers, JWT auth not yet shipped, no compliance attestations. An OSS launch does not need SOC 2, but it does need published benchmarks (even modest), shipped auth, and a security-posture doc — especially after the LiteLLM PyPI supply-chain compromise made the category security-paranoid. Turn that into an advantage: safety-first, in-process, no phone-home, deterministic, no LLM in the evaluation path.
- **License & governance:** Apache-2.0 is right. Decide the contribution model and whether to pursue a foundation path later (agentgateway's Linux Foundation move is the credibility benchmark).
- **OSS → monetization motion:** stars are the top of the funnel, not revenue (LiteLLM proves this). Have the enterprise-tier story ready at launch — SSO, audit, support — or risk adoption with no monetization path.

6.5 To verify before launch

- **Positioning vs. Bedrock Guardrails / AgentCore internally** — frame as complementary (content vs. action), not competitive with an AWS service.
- **"Gateway" vs. "guardrail" framing** — the guardrails market grows faster (~32.5% vs ~29.4% CAGR) and is where the M&A money is. Leaning "agent action guardrail" rather than "AI gateway" may place XYZ-guard in the higher-value, less-crowded category.

6.6 Bottom line

Own the action-safety niche — per-call risk score + human-in-the-loop intervene tier + behavioral safety — aimed at security/governance buyers protecting high-blast-radius agents, with the build-your-own crowd as the beachhead. Treat routing/cost/MCP as proof of completeness, not the pitch. Close the launch-credibility gaps (benchmarks, auth, security posture, name collision) before the repo goes public.

7. Caveats

- No analyst tracks vendor-level share for AI gateways; all "share" here is proxy-based.
- GitHub stars measure interest, not revenue or production deployments.
- Market-size figures come from single-vendor research reports and vary by how each defines "gateway"; the \$2.71B "inference gateways" figure partly includes edge/hardware inference, so it is an upper envelope, not a pure software-gateway TAM.

- Competitor performance numbers (Bifrost, agentgateway) are vendor self-benchmarked unless independently verified.

Sources: XYZ-guard repository (README, PRFAQ, CHANGELOG, gateway-architecture, control-plane docs); prior competitive research tracks; external market data (Business Research Company, Gartner category overview, Maxim, openalternative, Portkey, Deepak Gupta AI-gateway comparison).